

ASSIGNMENTS - 4

Assignment Date	17 October 2022
Student Name	DIVYA V
Student Roll Number	210519205013
Maximum Marks	2 Marks

Question:

Problem Statement: Abalone Age Prediction

Building a Regression Model

1. Download the dataset: Dataset
2. Load the dataset into the tool.
3. Perform Below Visualizations. · Univariate Analysis · Bi-Variate Analysis
Multi-Variate Analysis
4. Perform descriptive statistics on the dataset.
5. Check for Missing values and deal with them.
6. Find the outliers and replace them outliers
7. Check for Categorical columns and perform encoding.
8. Split the data into dependent and independent variables.
9. Scale the independent variables
10. Split the data into training and testing
11. Build the Model
12. Train the Model
13. Test the Model
14. Measure the performance using Metrics

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sb
```

```
In [ ]: ak=pd.read_csv("/content/abalone.csv")
```

```
In [ ]: ak.info
```

```
Out[32]: <bound method DataFrame.info of      Sex  Length  Diameter  Height  Whole weigh
t Shucked weight \
0      M    0.455    0.365    0.095      0.5140      0.2245
1      M    0.350    0.265    0.090      0.2255      0.0995
2      F    0.530    0.420    0.135      0.6770      0.2565
3      M    0.440    0.365    0.125      0.5160      0.2155
4      I    0.330    0.255    0.080      0.2050      0.0895
... ..
4172   F    0.565    0.450    0.165      0.8870      0.3700
4173   M    0.590    0.440    0.135      0.9660      0.4390
4174   M    0.600    0.475    0.205      1.1760      0.5255
4175   F    0.625    0.485    0.150      1.0945      0.5310
4176   M    0.710    0.555    0.195      1.9485      0.9455

      Viscera weight  Shell weight  Rings
0                0.1010      0.1500     15
1                0.0485      0.0700      7
2                0.1415      0.2100      9
3                0.1140      0.1550     10
4                0.0395      0.0550      7
... ..
4172              0.2390      0.2490     11
4173              0.2145      0.2605     10
4174              0.2875      0.3080      9
4175              0.2610      0.2960     10
4176              0.3765      0.4950     12

[4177 rows x 9 columns]>
```

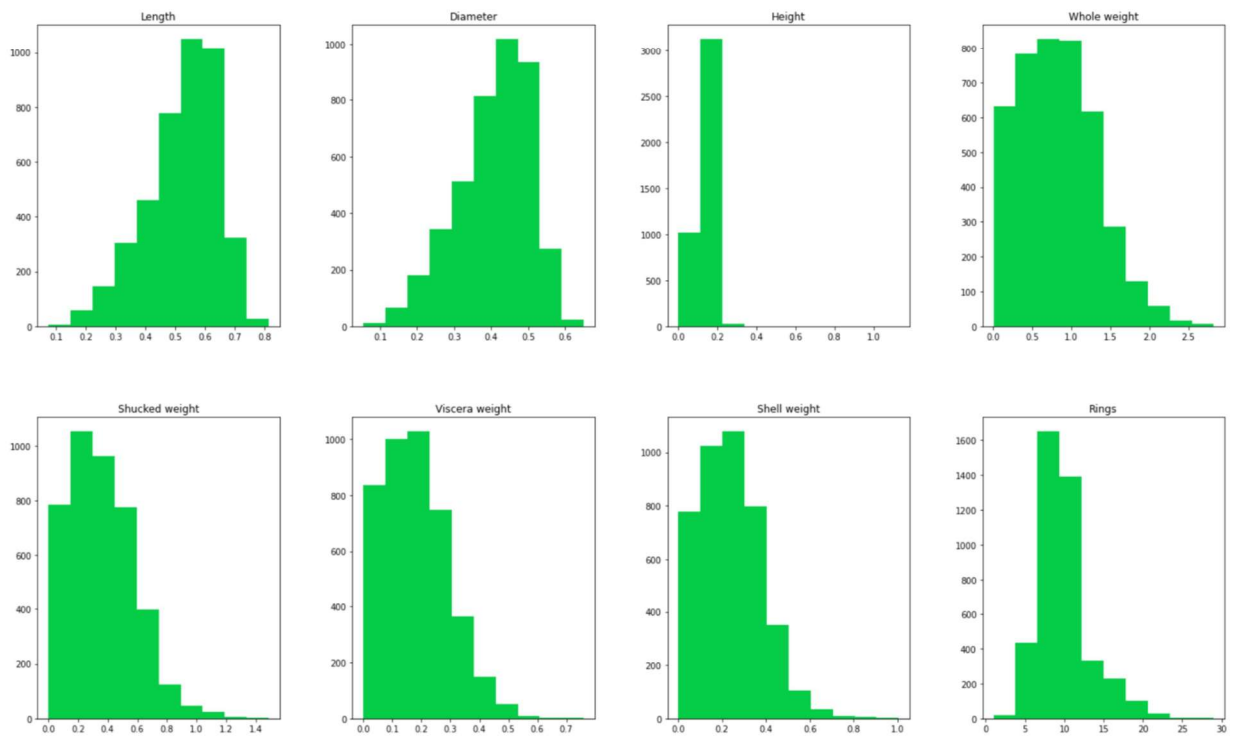
```
In [ ]: ak.isnull().sum()
```

```
Out[33]: Sex                0
Length                0
Diameter              0
Height                0
Whole weight          0
Shucked weight        0
Viscera weight        0
Shell weight          0
Rings                 0
dtype: int64
```

UNI-VARIATE ANALYSIS

```
In [ ]: ak.hist(figsize=(25,15),grid=False , color="#05cc46" , layout=(2,4))
```

```
Out[34]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f98071de550>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f98072cc1d0>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f98070ce510>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f9807084b10>],  
  [<matplotlib.axes._subplots.AxesSubplot object at 0x7f9807046150>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f980707d750>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f9807034dd0>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f9806ff8350>]],  
  dtype=object)
```



```
In [ ]: numerical_features=ak.select_dtypes(include=[np.number]).columns  
categorical_features = ak.select_dtypes(include=[np.object]).columns
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.

Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations> (<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>)

```
In [ ]: print(numerical_features)
categorical_features
```

```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight', 'Rings'],
      dtype='object')
```

```
Out[36]: Index(['Sex'], dtype='object')
```

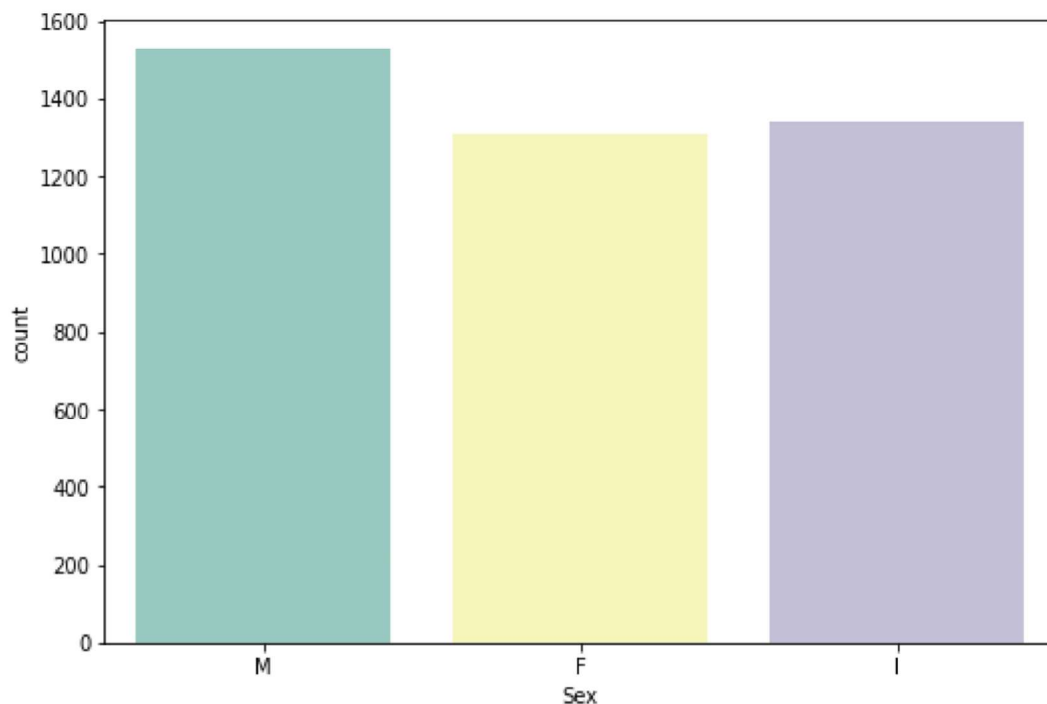
```
In [ ]: ak['age']=ak['Rings']+1.5
```

```
In [ ]: ak.age
```

```
Out[38]: 0      16.5
1       8.5
2      10.5
3      11.5
4       8.5
...
4172    12.5
4173    11.5
4174    10.5
4175    11.5
4176    13.5
Name: age, Length: 4177, dtype: float64
```

```
In [ ]: sb.countplot(x="Sex" , data=ak , palette='Set3')
```


```
Out[39]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9806cc67d0>
```



BI-VARIATE ANALYSIS

```
In [ ]: plt.figure(figsize = (20,7))
sb.heatmap(ak[numerical_features].corr(),annot=True)
```

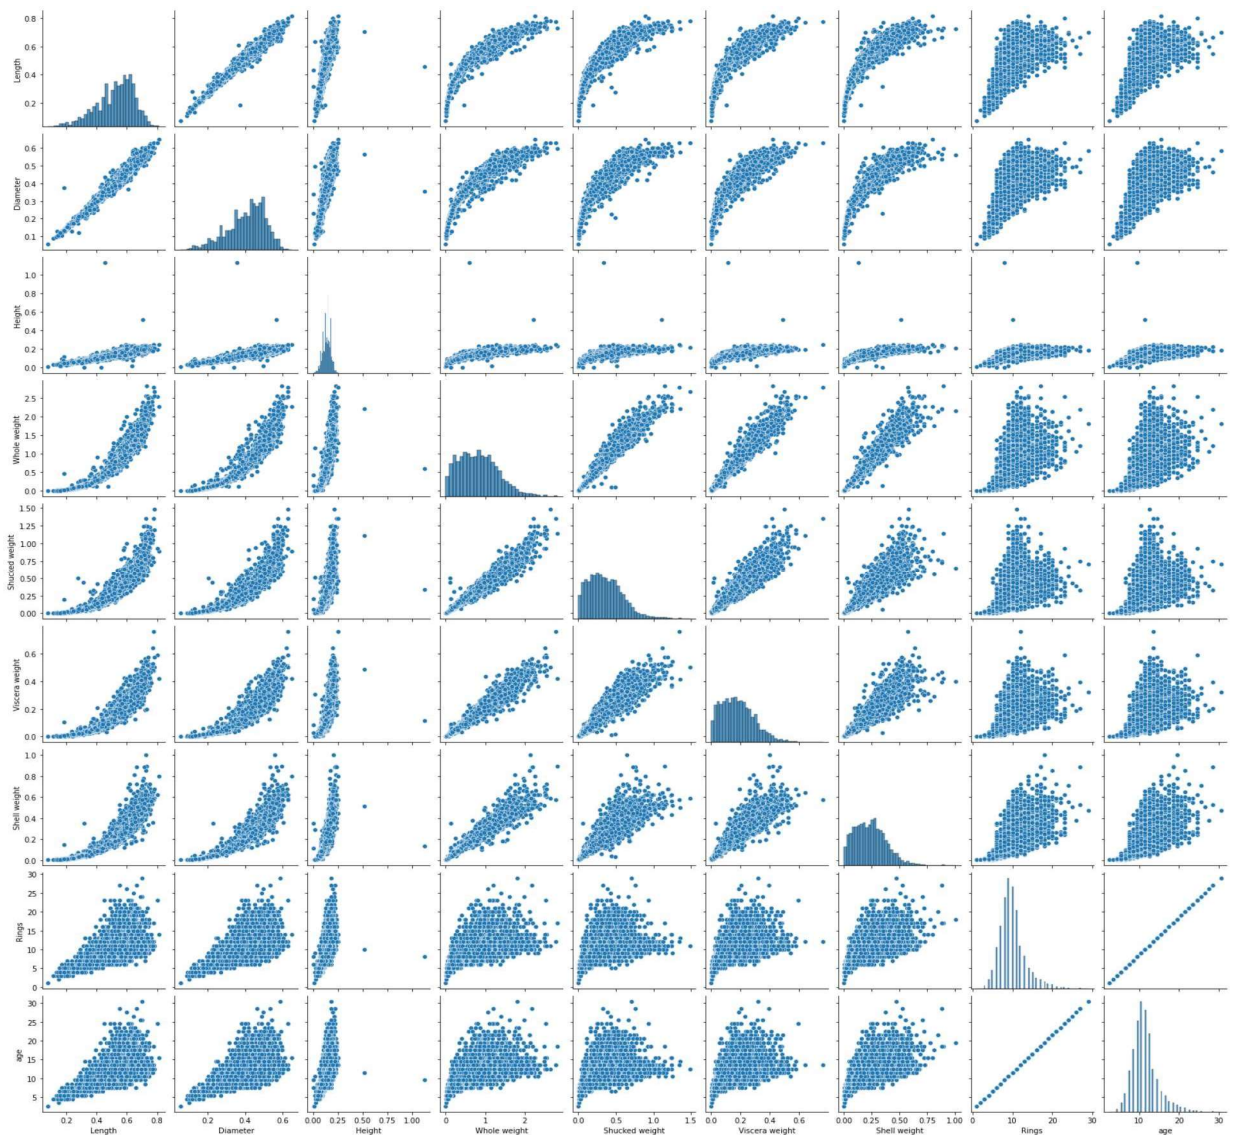
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x7f98097d8e90>



MULTI-VARIAT ANALYSIS

```
In [ ]: sb.pairplot(ak)
```

Out[41]: <seaborn.axisgrid.PairGrid at 0x7f98083d8c50>



Descriptive Statistics

```
In [ ]: ak.describe
```

```
Out[42]: <bound method NDFrame.describe of
ght  Shucked weight \
0      M    0.455    0.365    0.095    0.5140    0.2245
1      M    0.350    0.265    0.090    0.2255    0.0995
2      F    0.530    0.420    0.135    0.6770    0.2565
3      M    0.440    0.365    0.125    0.5160    0.2155
4      I    0.330    0.255    0.080    0.2050    0.0895
... ..
4172   F    0.565    0.450    0.165    0.8870    0.3700
4173   M    0.590    0.440    0.135    0.9660    0.4390
4174   M    0.600    0.475    0.205    1.1760    0.5255
4175   F    0.625    0.485    0.150    1.0945    0.5310
4176   M    0.710    0.555    0.195    1.9485    0.9455
```

```
Viscera weight  Shell weight  Rings  age
0              0.1010      0.1500    15  16.5
1              0.0485      0.0700     7   8.5
2              0.1415      0.2100     9  10.5
3              0.1140      0.1550    10  11.5
4              0.0395      0.0550     7   8.5
... ..
4172           0.2390      0.2490    11  12.5
4173           0.2145      0.2605    10  11.5
4174           0.2875      0.3080     9  10.5
4175           0.2610      0.2960    10  11.5
4176           0.3765      0.4950    12  13.5
```

```
[4177 rows x 10 columns]>
```

MISSING VALUES

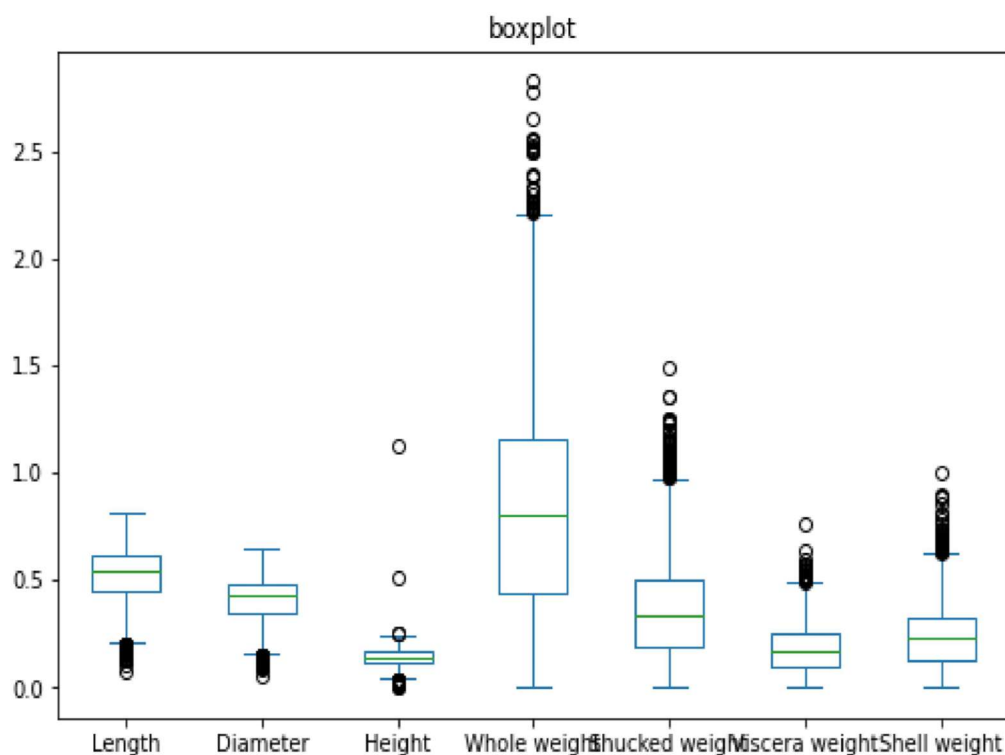
```
In [ ]: Missing_Values=ak.isnull().sum()
```

```
In [ ]: Missing_Values
```

```
Out[46]: Sex              0
Length              0
Diameter            0
Height              0
Whole weight        0
Shucked weight      0
Viscera weight      0
Shell weight        0
Rings               0
age                 0
dtype: int64
```

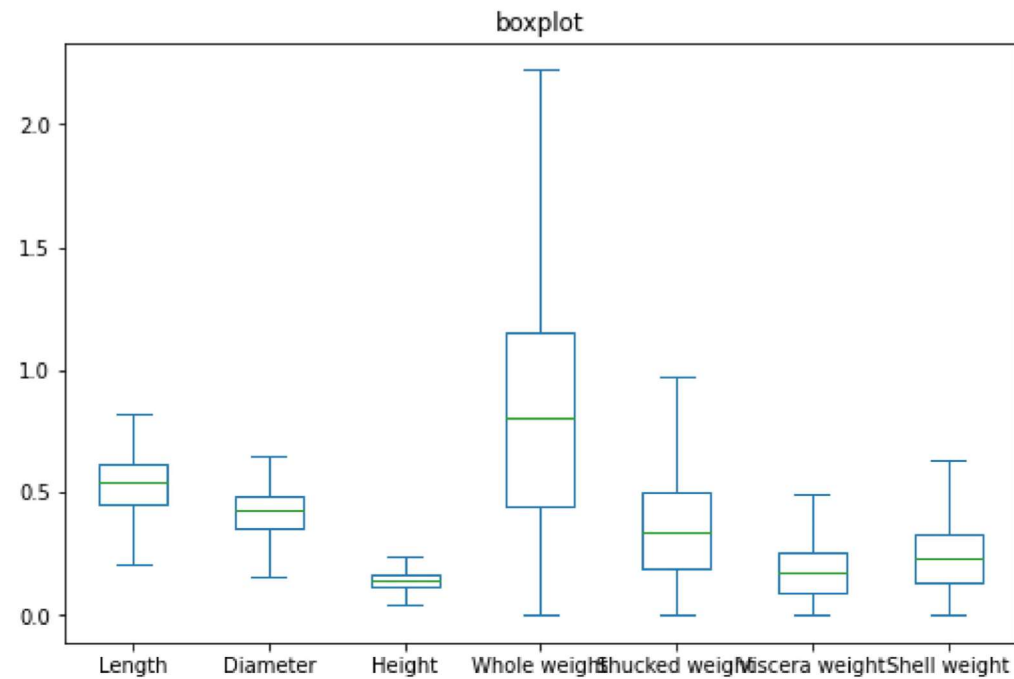

FINDING OUTLIERS AND REPLACING THEM

```
In [ ]: plt.rcParams['figure.figsize']=[8.50,5.50]  
ax = ak[['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight', 'Viscera we
```



```
In [ ]: for i in range(1,8):  
    Q1 = ak.iloc[:,i].quantile(0.25)  
    Q3 = ak.iloc[:,i].quantile(0.75)  
    IQR = Q3 - Q1  
    whisker_width = 1.5  
    lower_whisker = Q1 - (whisker_width*IQR)  
    upper_whisker = Q3 + (whisker_width*IQR)  
    ak.iloc[:,i] = np.where(ak.iloc[:,i]>upper_whisker,upper_whisker,np.where(ak.
```

```
In [ ]: plt.rcParams["figure.figsize"] = [8.50, 5.50]
ay = ak[['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight', 'Viscera we
```



Checking for Categorical columns and performing encoding

```
In [ ]: cate_data = ak.select_dtypes(include=['object']).copy()
```

```
In [ ]: from sklearn.preprocessing import LabelBinarizer

lb = LabelBinarizer()
lb_results = lb.fit_transform(cate_data['Sex'])
lb_results_df = pd.DataFrame(lb_results, columns=lb.classes_)

print(lb_results_df.head())
```

```
   F  I  M
0  0  0  1
1  0  0  1
2  1  0  0
3  0  0  1
4  0  1  0
```

```
In [ ]: result_df = pd.concat([cate_data, lb_results_df], axis=1)

print(result_df.head())
```

```
   Sex  F  I  M
0  M  0  0  1
1  M  0  0  1
2  F  1  0  0
3  M  0  0  1
4  I  0  1  0
```


Splitting into dependent and independent variables

```
In [ ]: x=ak.iloc[:,1:2].values  
x
```

```
Out[64]: array([[0.455],  
               [0.35 ],  
               [0.53 ],  
               ...,  
               [0.6   ],  
               [0.625],  
               [0.71 ]])
```

```
In [ ]: y=ak.age  
y
```

```
Out[54]: 0      16.5  
        1       8.5  
        2     10.5  
        3     11.5  
        4       8.5  
        ...  
       4172    12.5  
       4173    11.5  
       4174    10.5  
       4175    11.5  
       4176    13.5  
Name: age, Length: 4177, dtype: float64
```

```
In [ ]: x.shape
```

```
Out[55]: (4177, 7)
```

```
In [ ]: y.shape
```

```
Out[56]: (4177,)
```

Scaling the independent variables

```
In [ ]: print ("\n ORIGINAL VALUES: \n\n", x,y)
```

ORIGINAL VALUES:

```
[[0.455]  
 [0.35 ]  
 [0.53 ]  
 ...  
 [0.6   ]  
 [0.625]  
 [0.71 ]] 0      16.5  
1       8.5  
2     10.5  
3     11.5  
4       8.5  
...  
4172    12.5  
4173    11.5  
4174    10.5  
4175    11.5  
4176    13.5
```

Name: age, Length: 4177, dtype: float64

```
In [ ]: from sklearn import preprocessing
min_max_scaler = preprocessing.MinMaxScaler(feature_range =(0, 1))
new_y= min_max_scaler.fit_transform(x,y)
print ("\n VALUES AFTER MIN MAX SCALING: \n\n", new_y)
```

VALUES AFTER MIN MAX SCALING:

```
[[0.4122449 ]
 [0.24081633]
 [0.53469388]
 ...
 [0.64897959]
 [0.68979592]
 [0.82857143]]
```

SPLITTING THE DATA

```
In [ ]: X = ak.drop('age', axis = 1)
y = ak['age']
```

```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
x_train, x_test, y_train, y_test = train_test_split(new_y,x,test_size=0.2)
```

```
In [ ]: x_train.shape,x_test.shape
```

Out[78]: ((3341, 1), (836, 1))

```
In [ ]: y_train.shape,y_test.shape
```

Out[79]: ((3341, 1), (836, 1))

Building and training the model

```
In [ ]: model=LinearRegression()
model.fit(x_train,y_train)
```

Out[80]: LinearRegression()

Testing the model

```
In [ ]: y_pred=model.predict(x_test)
```

Measure the performance using Metrics

```
In [ ]: print('R Squared value:', r2_score(y_test,y_pred))
print('Mean Absolute Error:', mean_absolute_error(y_test,y_pred))
print('Mean Squared Error:', mean_squared_error(y_test,y_pred))
print('Root Mean Squared Error:', np.sqrt(mean_squared_error(y_test,y_pred)))
```

```
R Squared value: 1.0
Mean Absolute Error: 2.905039313836758e-17
Mean Squared Error: 2.853878008501086e-33
Root Mean Squared Error: 5.342169979045113e-17
```